

Vector Function Application Binary Interface Specification for AArch64

2021Q1

Date of Issue: 12th April 2021

arm

1 Preamble

1.1 Abstract

The *Vector Function Application Binary Interface Specification for AArch64* describes the application binary interface for vector functions generated by a compiler.

This document uses the [OpenMP](#) `declare simd` directive to classify the vector functions that can be associated to a scalar function.

The following rules apply to a compiler that implements the OpenMP directive `#pragma omp declare simd`:

1. The use of a `#pragma omp declare simd` construct for a *function definition* enables the creation of vector versions of the function from the scalar version of the function, that can be used to process multiple instances concurrently in a single invocation in a vector context (e.g. vectorized loops).
2. The use of a `#pragma omp declare simd` construct for a *function declaration* enables the compiler to know the exact list of available vector function implementations provided by a library that is based on the OpenMP pragmas found in the function's prototype of the library headers.

This Vector Function ABI defines a set of rules that the caller and the callee functions must obey.

The Vector Function ABI also describes how to use the `declare variant` directive introduced in OpenMP 5.0 to interface user-defined vector functions with a compiler.

1.2 Keywords

SVE

Scalable Vector Extension

A64

Instruction set of the ARMv8-A architecture

AArch64

64-bit execution mode of the ARMv8-A architecture

Advanced SIMD

SIMD and floating point instructions of the A64 instruction set

Qn register

Quad-word (128-bit) floating point register

Dn register

Double-word (64-bit) floating point register

Sn register

Single-word (32-bit) floating point register

Hn register

Half-word (16-bit) floating point register

Bn register

Byte (8-bit) floating point register

Vn register

Advanced SIMD vector register

Zn register

SVE vector register

ACLE

ARM C Language Extensions

SVE ACLE

ARM C Language Extensions for SVE

Vector function

A function processing vector data through the SIMD registers

Leaf function

Function at the end of a call tree

AAPCS

ARM Architecture Procedure Call Standard

AAELF64

ELF for the Arm 64-bit Architecture

OpenMP

[Open Multi-Processing standard](#)

Uniform parameter

A function parameter marked with the OpenMP *uniform* clause

Linear parameter

A function parameter marked with the OpenMP *linear* clause

LP64

Data model in which Long and Pointers are 64-bit

ILP32

Data model in which Integer, Long and Pointers are 32-bit

VLA

Vector Length Agnostic

VLS

Vector Length Specific

MTV (P)

P Maps To Vector

PBV (P)

P is Passed By Value

LS (P)

Lane Size of P

MAP (P)

Mapping of P

ADVSIMD_MAP (P)

Mapping of P - Advanced SIMD specific rules.

SVE_MAP (P)

Mapping of P - SVE specific rules.

NDS (f)

Narrowest Data Size of f

WDS (f)

Widest Data Size of f

1.3 Latest release and defects report

Please check [Application Binary Interface for the Arm® Architecture](#) for the latest release of this document.

Please report defects in this specification to the [issue tracker page on GitHub](#).

1.4 License

This work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/4.0/> or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.

Grant of Patent License. Subject to the terms and conditions of this license (both the Public License and this Patent License), each Licensor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable (except as stated in this section) patent license to make, have made, use, offer to sell, sell, import, and otherwise transfer the Licensed Material, where such license applies only to those patent claims licensable by such Licensor that are necessarily infringed by their contribution(s) alone or by combination of their contribution(s) with the Licensed Material to which such contribution(s) was submitted. If You institute patent litigation against any entity (including a cross-claim or counterclaim in a lawsuit) alleging that the Licensed Material or a contribution incorporated within the Licensed Material constitutes direct or contributory patent infringement, then any licenses granted to You under this license for that Licensed Material shall terminate as of the date such litigation is filed.

1.5 About the license

As identified more fully in the [License](#) section, this project is licensed under CC-BY-SA-4.0 along with an additional patent license. The language in the additional patent license is largely identical to that in Apache-2.0 (specifically, Section 3 of Apache-2.0 as reflected at <https://www.apache.org/licenses/LICENSE-2.0>) with two exceptions.

First, several changes were made related to the defined terms so as to reflect the fact that such defined terms need to align with the terminology in CC-BY-SA-4.0 rather than Apache-2.0 (e.g., changing “Work” to “Licensed Material”).

Second, the defensive termination clause was changed such that the scope of defensive termination applies to “any licenses granted to You” (rather than “any patent licenses granted to You”). This change is intended to help maintain a healthy ecosystem by providing additional protection to the community against patent litigation claims.

1.6 Contributions

Contributions to this project are licensed under an inbound=outbound model such that any such contributions are licensed by the contributor under the same terms as those in the LICENSE file.

1.7 Trademark notice

The text of and illustrations in this document are licensed by Arm under a Creative Commons Attribution–Share Alike 4.0 International license (“CC-BY-SA-4.0”), with an additional clause on patents. The Arm trademarks featured here are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. Please visit <https://www.arm.com/company/policies/trademarks> for more information about Arm’s trademarks.

1.8 Copyright

Copyright (c) 2018-2021, Arm Limited and its affiliates. All rights reserved.

Contents

1	Preamble	2
1.1	Abstract	2
1.2	Keywords	2
1.3	Latest release and defects report	4
1.4	License	5
1.5	About the license	5
1.6	Contributions	5
1.7	Trademark notice	5
1.8	Copyright	5
2	About this document	8
2.1	Change control	8
2.1.1	Current status and anticipated changes	8
2.1.2	Change history	8
3	Definitions	10
3.1	Vector Procedure Call Standard	10
3.1.1	AAVPCS Table	10
3.2	Dynamic linking for AAVPCS	11
3.3	Extended Vector Notation for Advanced SIMD	11
3.3.1	Padded Short Vectors	11
3.3.2	Extended short vectors	12
3.4	SVE unpacked vector	13
3.5	Complex types	13
4	Vector function signature	14
4.1	Common rules for parameter mapping	14
4.1.1	Maps To Vector	14
4.1.2	Pass By Value	14
4.1.3	Parameter and return value mapping	15
4.2	Vector length selection	15
4.2.1	Lane Size of a function parameter / return value	15
4.2.2	Narrowest and Widest Data Size of a Function	15
4.3	Advanced SIMD-specific rules	16
4.3.1	Vector Length (Advanced SIMD)	16
4.3.2	Parameter mapping (Advanced SIMD)	16
4.4	SVE-specific rules	16
4.4.1	Vector Length (SVE)	16
4.4.2	Parameter mapping (SVE)	16
4.4.3	Unpacked parameters / return value	17

4.5	Vector function name mangling	17
4.5.1	Name mangling function	17
4.5.2	Description of the parameter_kind token	18
4.6	Advanced SIMD examples	19
4.7	SVE Examples	20
4.8	Linear parameters examples	21
4.8.1	Linear clause for integral parameters	21
4.8.2	Linear clause for pointer parameters	21
4.8.3	Linear clause for integral reference parameters	22
5	User defined vector functions	24
5.1	AArch64 Variant Traits	24
5.2	Examples	25
6	Masking	27
6.1	Inactive lanes	27
6.2	Advanced SIMD masking	27
6.2.1	Example on Complex Masking	28
6.3	SVE masking	29
6.3.1	Generating the predicate value of the mask parameter	29
7	Additional examples	30
7.1	RGB Example	31
8	Footnotes	33

2 About this document

2.1 Change control

2.1.1 Current status and anticipated changes

The following support level definitions are used by the Arm ABI specifications:

Release

Arm considers this specification to have enough implementations, which have received sufficient testing, to verify that it is correct. The details of these criteria are dependent on the scale and complexity of the change over previous versions: small, simple changes might only require one implementation, but more complex changes require multiple independent implementations, which have been rigorously tested for cross-compatibility. Arm anticipates that future changes to this specification will be limited to typographical corrections, clarifications and compatible extensions.

Beta

Arm considers this specification to be complete, but existing implementations do not meet the requirements for confidence in its release quality. Arm may need to make incompatible changes if issues emerge from its implementation.

Alpha

The content of this specification is a draft, and Arm considers the likelihood of future incompatible changes to be significant.

Unless otherwise indicated, all content in this document is at the **Release** quality level.

2.1.2 Change history

If there is no entry in the change history table for a release, there are no changes to the content of the document for that release.

Issue	Date	Change
2Q2018	26th June 2018	First public release.
2019Q1	29th March 2019	Fix broken link in License section. Fix parameter numbering for linear steps in Vector function name mangling . Clarify the behavior for structures like <code>struct { int8_t R, G, B; };</code> in Parameter and return value mapping , and relative RGB Example .
2019Q1.1	30th April 2019	Minor clarification on the definition of SVE unpacked vector . Refer to the original AAPCS and list the registers that are call-preserved and call-clobbered in the base convention (Vector Procedure Call Standard , no functional change). Add chapter on User defined vector functions via OpenMP 5.0.

Issue	Date	Change
2019Q2	30th June 2019	<p>Fix the use of <code>declare variant</code> in User defined vector functions via OpenMP 5.0.</p> <p>Add section on Dynamic linking for AAVPCS with new requirement for ELF platforms that support dynamic linking.</p> <p>Fix mangled name for function <code>bar</code> in Example on Complex Masking.</p> <p>Non functional changes:</p> <ol style="list-style-type: none"> 1. Split the table on integral value and pointers in the Linear parameters examples into two separate tables, Linear clause for integral parameters and Linear clause for pointer parameters. 2. Extend the information of Linear clause for integral parameters, Linear clause for pointer parameters and Linear clause for integral reference parameters in the section on the Linear parameters examples, to include the mapping to the token of the mangled name and specify the cases in which the size of the underlying data type must be used as multiplier for the <code>step</code>. 3. In the section on the Vector function name mangling, change the type of numbers used in the token of the linear parameters from decimal to <code>integrals</code>, and improve the description of the rules.
2019Q4	30th January 2020	<p>Github preview release with an open source license.</p> <p>Major changes:</p> <ol style="list-style-type: none"> 1. New License, with relative explanation in About the license. 2. New sections on Contributions, Trademark notice, and Copyright. <p>Minor changes:</p> <ol style="list-style-type: none"> 1. The explanation of RGB Example has gained item 5, that refers to the rule that renders the return value as the first input parameter. <p>Several changes have been applied to the sources to fix the rendered page produced by github.</p> <p>In particular:</p> <ol style="list-style-type: none"> 1. The following sections have been renamed to make the implicit link associated to them unique: Vector Length (Advanced SIMD), Parameter Mapping (Advanced SIMD), Vector Length (SVE), Parameter Mapping (SVE) 2. The following sections have been added to be able to cross-reference tables and code examples that cannot be referenced using standard rST markup: AAVPCS Table, Name mangling function, Linear clause for integral parameters, Linear clause for pointer parameters, Linear clause for integral reference parameters, AArch64 Variant Traits.
2020Q2	1st July 2020	<p>Clarify whether <code>aarch64_vector_pcs</code> is needed for SVE.</p> <p>Clarify the definition of complex type.</p>

3 Definitions

3.1 Vector Procedure Call Standard

AArch64 functions use the calling convention described in section 5 of the *Procedure Call Standard for the ARM 64-bit Architecture (with SVE support)*, or AAPCS hereafter. The most recent version of the AAPCS can be found on developer.arm.com.

Note

The SVE-specific rules of the AAPCS are in beta version. The list of SVE call-clobbered and call-preserved registers in table [AAVPCS Table](#) will be updated when the final version of the AAPCS is published.

The procedural calling standard of the AAPCS requires that none of the 32 Advanced SIMD vector registers V0-V31 are treated as call-preserved (with the exception of the lower half of V8-V15, or D8-D15), thus requiring the caller to perform up to 32 vector stores before a call and up to 32 vector loads after it (see section 5.1.2 of AAPCS). For workloads with performance hot spots in leaf routines (an example of which are vector math functions), we find that a modified procedural calling standard for the vector units in AArch64 would be more efficient than the base procedural calling standard. Therefore, to efficiently support such vector routines, we define a modified version of the base procedural calling standard, called the *Vector Procedure Call Standard for the Arm 64-bit Architecture (AAVPCS)*.

The list of parameter, result, call-preserved and call-clobbered registers for the AAVPCS are presented in the following table:

3.1.1 AAVPCS Table

Modified PCS for vector functions (AAVPCS)

Extension	Parameter and Result registers	Call-clobbered registers	Call-preserved registers
Advanced SIMD	V0-V7	V0-V7, V24-V31	V8-V23
SVE	Z0-Z7	See AAPCS	

The AAVPCS is implicit when a `#pragma omp declare simd` clause is attached to a function definition or declaration. For user-defined Advanced SIMD vector functions, the same behavior can be obtained by adding the `aarch64_vector_pcs` function attribute to the function definition or declaration as in the following examples. For user-defined SVE vector functions the attribute is not required as AAPCS and AAVPCS are equivalent. Note that to ensure the compiler produces ABI consistent code, the attribute must be specified in every declaration and definition of the function.

```
/* function definition */
__attribute__((aarch64_vector_pcs))
uint64x2_t foo(uint32x2_t a, float32x2_t b) {
    /* function body */
}
/* function declaration */
__attribute__((aarch64_vector_pcs)) float64x2_t bar(float64x2_t a);
```

3.2 Dynamic linking for AAVPCS

On ELF platforms with dynamic linking support, symbol definitions and references must be marked with the `STO_AARCH64_VARIANT_PCS` flag set in their `st_other` field if the following conditions hold:

1. The binding for the symbol is not `STB_LOCAL`, or it is in the dynamic symbol table.
2. The symbol is associated with a function following the AAVPCS convention.

For more information on `STO_AARCH64_VARIANT_PCS`, see [AAELF64](#).

Note

Marking all functions that follow the AAVPCS convention is a valid way of implementing this requirement.

3.3 Extended Vector Notation for Advanced SIMD

For the purposes of this specification, we define the following notational extensions for the Advanced SIMD vector types defined by the AAPCS64. These types are not made available to the user.

3.3.1 Padded Short Vectors

Padded short vectors extend the definition of short vectors and are used as a notational convenience to describe vector types with a size of less than 64 bits. These can be formed where the `simdlen` clause specified in an OpenMP `declare simd` construct would force a smaller vector than would meet the AAPCS definition of a short vector. These have the form of a vector with `<N>` elements of type `<T>`:

```
<T>x<N>_t
```

Where

```
sizeof(<T>) * <N> < 8
```

A padded short vector is represented as an 8-byte short vector type with elements of type `<T>` in which lanes `<N>` and above have unspecified values. For example, a padded short vector `uint16x2_t` is represented as a `uint16x4_t` in which lanes 2 and 3 have unspecified values.

The contents of the 8-byte vector are arranged as though the whole padded short vector were a single lane. For example, a `uint16x2_t` is stored in the `uint16x4_t` as though it were lane 0 in a `uint32x2_t`.

Note

When a padded short vector is transferred between registers and memory it is treated as an opaque object of the notional type. That is, a padded short vector is stored in memory as if it were stored with a single STR of an object of the size of the notional type of the padded short vector; a padded short vector is loaded from memory using the corresponding LDR instruction. On a little-endian system this means that element 0 will always contain the lowest addressed element of a padded short vector; on a big-endian system element 0 will contain the highest-addressed element of a padded short vector.

This is shown in the following table.

Big-endian vs Little-endian padded vectors.

Padded short vector type	Short vector type	Little-endian	Big-endian
[u]int8x2_t	[u]int8x8_t	X X X X X X A[1] A[0]	X X X X X X A[0] A[1]
[u]int8x4_t	[u]int8x8_t	X X X X A[3] ... A[0]	X X X X A[0] ... A[3]
float16x2_t	float16x4_t	X X A[1] A[0]	X X A[0] A[1]

The set of padded short vector types, the short vector type they map to, and the appropriate store width for each type is given in the following table,

Padded short vectors

Padded short vector type	Short vector type	LDR/STR registers
[u]int8x1_t	[u]int8x8_t	Bn
[u]int8x2_t	[u]int8x8_t	Hn
[u]int8x4_t	[u]int8x8_t	Sn
[u]int16x1_t	[u]int16x4_t	Hn
[u]int16x2_t	[u]int16x4_t	Sn
float16x1_t	float16x4_t	Hn
float16x2_t	float16x4_t	Sn
float32x1_t	float32x2_t	Sn

When using a padded short vector, the contents of the elements of the associated short vector that lie outside the padded short vector are undefined.

Where padded short vectors are used, this may cause the compiler to emit conservative, scalar code to process their content.

No language bindings are provided for padded short vectors. Padded short vectors are not generated for declare simd constructs with no simdlen clause.

3.3.2 Extended short vectors

Extended short vectors extend the AAPCS definition of short vectors and are used as a notational convenience to describe vector types with a size greater than 128 bits. These can be formed where the required vectorization factor would create a larger vector than would meet the AAPCS definition of a short vector. These have the form:

```
<T>x<N>_t
```

Where

```
sizeof(<T>) * <N> > 16
```

Extended short vectors are represented as a structure containing an array of short vectors of the appropriate type. These have the general form:

```
struct <T>x<NN>x<M>_t { <T>x<NN>_t val[<M>]; };
```

Where $\langle NN \rangle$ is such that $\langle N \rangle = \langle NN \rangle * \langle M \rangle$.

A subset of the possible vector types are given in the following table.

Extended short vector examples

Notional type	Parameter/Return type
<code>int32x16_t</code>	<code>struct int32x4x4_t { int32x4_t val[4]; };</code>
<code>float64x4_t</code>	<code>struct float64x2x2_t { float64x2_t val[2]; };</code>
<code>int32x16_t</code>	<code>struct int32x4x4_t { int32x4_t val[4]; };</code>

No language bindings are provided for extended short vectors, though some of these types are also defined by `arm_neon.h`.

3.4 SVE unpacked vector

Let `sv<T>_t` be an SVE ACLE vector type with lanes of type `<T>`. The vector is said to be *unpacked* if only the logical lanes corresponding to the multiples of some power of 2 greater or equal than 2 can be set active by a `svbool_t` predicate. Conversely, the vector is said to be *packed* if any lane can be active.

For example, 32-bit signed integers from a reference `int32_t * A` can be loaded into an unpacked `svint32_t` vector at lanes 0, 2, 4,... and so on, effectively using only half of the lanes available in the vector. In the following example, the resulting SVE packed vector is shown together with two unpacked versions (X is for undefined content):

lane idx	8	7	6	5	4	3	2	1	0	
[msb]	...	A[7]	A[6]	A[5]	A[4]	A[3]	A[2]	A[1]	A[0]	[lsb] // packed
[msb]	...	X	A[3]	X	A[2]	X	A[1]	X	A[0]	[lsb] // unpacked 0, 2, 4, ...
[msb]	...	X	X	X	A[1]	X	X	X	A[0]	[lsb] // unpacked 0, 4, 8, ...

3.5 Complex types

In this specification, the term *complex type* will be used to refer to the following language-dependent types:

1. For C/C++, any of the complex types as defined in the C99 and C11 standards. For more information, see [C complex numbers](#).
2. For C++, any of the complex types as defined in the header file `<complex>`. For more information, see [C++ complex numbers](#).
3. For Fortran, any of the types defined using the `COMPLEX` statement.

4 Vector function signature

This section describes how the scalar functions decorated with the OpenMP `declare simd` pragma are associated to vector function signatures.

When vectorizing the following loop, whatever vectorization factor we choose, we want to make sure that the compiler expects a vector version of `f` and `g` that operates on the same number of lanes.

```
float f(double);
double g(float);
float x[];
//...
for (int i = 0; i < 100; ++i)
    x[i] = f(g(x[i]));
```

The rules given in this chapter guarantee that any `#pragma omp declare simd` attached to a *function declaration* or *definition* generates a unique set of vector functions associated to the original scalar function. This is done to make sure that library vendors can provide a unique way to interface the routines of the library with a compiler, by means of the `declare simd` directive.

In all cases, the order of the vector function parameters reflects the ordering of the parameters of the original scalar function.

Throughout this chapter, `f` is a function declaration or definition decorated with an OpenMP `declare simd` directive, `<P>` is the return value or an input parameter of `f`, and `<T(P)>` is its associated type.

4.1 Common rules for parameter mapping

One or more vector functions `F` are associated to the original scalar function `f`. The return value and each function parameter is mapped to a unique return value or input parameter respectively, named *Mapping of P*, or `MAP(P)`. The type of these vector function return value and input parameters depends on the following rules. Their order is the same as in the original scalar function `f`.

4.1.1 Maps To Vector

To each `<P>`, a true / false predicate "`P Maps To Vector`", or `MTV(P)` hereafter, is associated as follows:

1. If `<P>` is an input parameter such that:
 1. `<P>` is a uniform value, or
 2. `<P>` is a linear value and not a reference marked with `val` or no linear modifiers,then `MTV(P)` is false.
2. If `P` is a void return value, then `MTV(P)` is false;
3. In all other cases, `MTV(P)` is true.

4.1.2 Pass By Value

When a scalar parameter maps to a vector, that vector sometimes contains the *values* of the scalar parameters and sometimes contains the *addresses* of the scalar parameters. The predicate *Pass by Value* `PBV(T)` is true if the former case applies for scalar parameters of type `T`; it is false if the latter case applies. The predicate is defined as follows:

1. `PBV(T)` is true if (a) `T` is an integer, floating-point or pointer type and (b) `sizeof(T)` is 1, 2, 4 or 8.
2. `PBV(T)` is true if `T` is a complex type with components of type `T'` and if `PBV(T')` is true.

3. Otherwise $PBV(T)$ is false.

4.1.3 Parameter and return value mapping

When mapping the return value or an input parameter $\langle P \rangle$ of the scalar function to the corresponding $MAP(P)$ in the vector function, the following rules apply:

1. If $MTV(P)$ is false, then $MAP(P)$ is P .
2. Otherwise, if $MTV(P)$ is true, then $MAP(P)$ is target specific:
 1. For Advanced SIMD, $MAP(P) = ADVSIMD_MAP(P)$, with $ADVSIMD_MAP(P)$ defined in section [Advanced SIMD-specific rules](#).
 2. For SVE, $MAP(P) = SVE_MAP(P)$, with $SVE_MAP(P)$ defined in section [SVE-specific rules](#).
3. In all cases, when $\langle P \rangle$ is the return value, and:
 1. $MTV(P) = true$.
 2. $PBV(P) = false$.
 3. $MAP(P)$ is a vector of pointers.

Then the return type of the associated vector function is `void`, and $MAP(P)$ becomes the first parameter of the vector function. The caller is responsible for allocating the memory associated with the pointers in $MAP(P)$.

4.2 Vector length selection

A set of vector lengths $VLEN$ is sometimes associated with the generated vector function F . When this is done, the algorithm for selecting the value(s) of $VLEN$ is target dependent. The algorithm makes use of the definitions in this section.

4.2.1 Lane Size of a function parameter / return value

We then define the *Lane Size of P* , or $LS(P)$, as follows.

1. If $MTV(P)$ is false and P is a pointer or reference to some type T for which $PBV(T)$ is true, $LS(P) = sizeof(T)$.
2. If $PBV(T(P))$ is true, $LS(P) = sizeof(P)$.
3. Otherwise $LS(P) = sizeof(uintptr_t)$.

4.2.2 Narrowest and Widest Data Size of a Function

For the function f , we define the following concepts:

1. The *Narrowest Data Size of f* , or $NDS(f)$, as the minimum of the lane size $LS(P)$ among all input parameters and return value $\langle P \rangle$ of f .
2. The *Widest Data Size of f* , or $WDS(f)$, as the maximum of the lane size $LS(P)$ among all input parameters and return value $\langle P \rangle$ of f .

Note that by definition the value of $NDS(f)$ and $WDS(f)$ can only be 1, 2, 4, 8, and 16.

4.3 Advanced SIMD-specific rules

This section describes the Advanced SIMD-specific rules for mapping $\langle P \rangle$ to its corresponding vector parameter $\text{MAP}(P)$ when $\text{MTV}(P) = \text{true}$.

4.3.1 Vector Length (Advanced SIMD)

A VLEN is always associated with the vector function. The rules to generate the set of the available values are:

1. If $\text{simdlen}(\text{len})$ is specified, then the compiler generates only one version with $\text{VLEN} = \text{len}$. The value of vlen must be a power of 2.
2. If no simdlen is specified, the compiler generates multiple versions, according to the following rules:
 1. if $\text{NDS}(f) = 1$, then $\text{VLEN} = 16, 8$;
 2. if $\text{NDS}(f) = 2$, then $\text{VLEN} = 8, 4$;
 3. if $\text{NDS}(f) = 4$, then $\text{VLEN} = 4, 2$;
 4. if $\text{NDS}(f) = 8$ or $\text{NDS}(f) = 16$, then $\text{VLEN} = 2$.

4.3.2 Parameter mapping (Advanced SIMD)

For a value of VLEN , the $\text{ADVSIMD_MAP}(P)$ is build as follows:

1. If $\text{PBV}(T(P))$ is false, $\text{ADVSIMD_MAP}(P)$ is a vector of VLEN elements of type uintptr_t .
2. If $T(P)$ is a complex type with components of type T , $\text{MAP}(P)$ is a vector of $2 * \text{VLEN}$ elements of type T .
3. Otherwise $\text{ADVSIMD_MAP}(P)$ is a vector of VLEN elements of type $T(P)$.
4. An optional $\{\text{not}\}\text{inbrach}$ clause defines whether or not a vector mask parameter is added as the last input parameter of F , according to the rules in table 1 in chapter 4. The vector mask type is selected by building a vector of VLEN elements consisting of unsigned integers of $\text{NDS}(f)$ bytes. The generation of the values in the mask parameter is described in section 4.1.

4.4 SVE-specific rules

This section describes the SVE-specific rules for mapping $\langle P \rangle$ to its corresponding vector parameter $\text{MAP}(P)$ when $\text{MTV}(P) = \text{true}$.

One vector function F is associated to f depending on its classification via the `declare simd` directive. The vector signatures that get generated are the same in all cases.

4.4.1 Vector Length (SVE)

1. If no simdlen clause is specified, a VLA vector version is associated.
2. When using a $\text{simdlen}(\text{len})$ clause, the compiler expects a VLS vector version of the function that is tuned for a specific implementation of SVE. The size of the implementation is $\text{WDS}(f) * \text{len} * 8$.

4.4.2 Parameter mapping (SVE)

Whether targeting VLA SVE or VLS SVE, the rules for mapping $\langle P \rangle$ to $\text{SVE_MAP}(P)$ are:

1. If $\text{PBV}(T(P))$ is false, $\text{SVE_MAP}(P)$ is a scalable vector of uintptr_t .

2. If $T(P)$ is a complex type with components of type T , $SVE_MAP(P)$ is a scalable vector of T .
3. Otherwise $SVE_MAP(P)$ is a scalable vector of $T(P)$.
4. An additional `svbool_t` mask parameter is added as the last parameter of F . The generation of the mask values is described in section 4.2.

4.4.3 Unpacked parameters / return value

The vectors of the signature of F are packed or unpacked according to the following rules:

1. if $LS(P) = WDS(F)$, then the vector is packed.
2. If $LS(P) < WDS(F)$, then the vector is unpacked.

Each element in the unpacked vector occupies the same number of bits as in the packed vector, and all elements are aligned to their least significant bits.

The following example shows the contents of an SVE vector consisting of 1-byte lanes, unpacked and aligned with the 4-byte lanes of a packed vector. The ?? characters indicate a byte whose value is undefined.

```
Zn.b [msb] ... 0x??????03 0x??????02 0x??????01 0x??????00 [lsb]
Zn.s [msb] ... 0x00000003 0x00000002 0x00000001 0x00000000 [lsb]
```

4.5 Vector function name mangling

The rules of the mangling scheme for vector functions are summarized by [Name mangling function](#).

With reference to [Name mangling function](#), the rules for building the `<parameters>` group are:

1. We generate one `<parameter>` token in the `<parameters>` group for each of the input parameters of the scalar function. The tokens are in the same order as the input parameters.
2. The rules for choosing the `<parameter_kind>` are defined in the [Description of the parameter_kind token](#).
3. The optional "a" `<X>` token represents the alignment value (in bytes) specified in the `aligned` clause (for example `aligned(c:a)`).
 1. When targeting Advanced SIMD, if the value `a` is missing, the default alignment value is 16 (128 bits), so that an `aligned` clause with no alignment is mangled as `a16`.
 2. When targeting SVE, the default value of an `aligned` clause is the alignment of the type pointed to by the corresponding parameter of the scalar signature. For example, `aligned(x)` for `T *x` defaults to the value `_Alignof(typeof(T))`.

4.5.1 Name mangling function

Name mangling grammar for vector functions.

```
<vector name> := <prefix> "_" <name>

<name> := Assembly name of the function

<prefix> := "_ZGV" <isa> <mask> <len> <parameters>

<isa> := "n"    (Advanced SIMD)
        | "s"    (SVE)
```

```

<mask> := "N" (No Mask)
        | "M" (Mask)

<len> := VLEN (VLS SVE or Advanced SIMD)
        | "x" (VLA SVE)

<parameters> := <parameter> { <parameter> }

<parameter> := <parameter_kind> [ "a" <X> ]

                                OpenMP version
                                support (onwards)
<parameter_kind> := "v"                4.0
                    | "l" | "l" <number> 4.0
                    | "R" | "R" <number> 4.5
                    | "L" | "L" <number> 4.5
                    | "U" | "U" <number> 4.5
                    | "ls" <pos>         4.5
                    | "Rs" <pos>         4.5
                    | "Ls" <pos>         4.5
                    | "Us" <pos>         4.5
                    | "u"                4.0

<number> := "n" <X> // "n" means negative
           | <Y>

<pos> := <X>

<X> := integral number greater than or equal to 1
<Y> := integral number greater than or equal to 2

```

4.5.2 Description of the *parameter_kind* token

4.5.2.1 No clause

"v"

Vector parameter - default for no linear/uniform clause.

4.5.2.2 uniform *clause*

"u"

Uniform parameter specified in the uniform clause. For example, `uniform(c)`.

4.5.2.3 linear *clause when* step *is a compile time constant*

"l" | "l" <number>

Linear parameter <P> for which (a) the step is a compile-time constant, (b) `MTV(P)=false` and (c) the linear clause has either a *val* modifier or no modifier. <number> is the value of the constant linear step, or an empty string if the step is 1. For example, `linear(i:2)` gives 12 and `linear(i:1)` gives 1 when the type of *i* is integer.

"R" | "R" <number>

Linear parameter <P> for which (a) the step is a compile-time constant, and (b) the linear clause has a *ref* modifier. <number> is the value of the constant linear step, or an empty string if the step is 1. For example, `linear(ref(i):3)` gives R3 and `linear(ref(i):1)` gives R when the type of *i* is integer.

"L" | "L" <number>

Linear parameter `<P>` for which (a) the step is a compile-time constant, (b) `MTV(P)=true` and (c) the linear clause has either a `val` modifier or no modifier. `<number>` is the value of the constant linear step, or an empty string if the step is 1. For example, `linear(val(i):-3)` gives `Ln3` when the type of `i` is integer.

In the previous cases, when the parameter `<P>` marked by the linear clause is a pointer or an OpenMP integral reference to a type `T`, the step of the linear clause must be multiplied by the size in bytes of the pointee, so that `<number>=sizeof(T) x step`.

"U" | "U" `<number>`

Linear parameter `<P>` for which (a) the step is a compile-time constant and (b) the linear clause has a `uval` modifier. `<number>` is the value of the constant linear step, or an empty string if the step is 1. For example, `linear(uval(i):2)` gives `U2`.

4.5.2.4 linear clause when step is a loop-independent runtime invariant

"ls" `<pos>`

Linear parameter `<P>` for which (a) the step is a loop-independent runtime invariant, (b) `MTV(P)=false` and (c) the linear clause has either a `val` modifier or no modifier. `<pos>` is the position (starting from 0) of the step parameter specified in the uniform clause (required by the OpenMP specs). For example, `linear(i:c) uniform(c)` with `c` being the third parameter gives `ls2`.

"Rs" `<pos>`

Linear parameter `<P>` for which (a) the step is a loop-independent runtime invariant and (b) the linear clause has a `ref` modifier. `<pos>` is the position of the step parameter (starting from 0) specified in the uniform clause (required by the OpenMP specs). For example, `linear(ref(i):c) uniform(c)` with `c` being the third parameter gives `Rs2`.

"Ls" `<pos>`

Linear parameter `<P>` for which (a) the step is a loop-independent runtime invariant, (b) `MTV(P)=true` and (c) the linear clause has either a `val` modifier or no modifier. `<pos>` is the position of the step parameter (starting from 0) specified in the uniform clause (required by the OpenMP specs). For example, `linear(val(i):c) uniform(c)` with `c` being the first parameter, gives `Ls0`.

"Us" `<pos>`

Linear parameter `<P>` for which (a) the step is a loop-independent runtime invariant and (b) the linear clause has a `uval` modifier. `<pos>` is the position of the step parameter (starting from 0) specified in the uniform clause (required by the OpenMP specs). For example, `linear(uval(i):c) uniform(c)` with `c` being the third parameter, gives `Us2`.

4.6 Advanced SIMD examples

The following example shows which vector versions are provided when no `simdlen` clause is attached to the `declare simd` directive of a *function declaration*.

```
#pragma omp declare simd
float f(double x);
#pragma omp declare simd
double g(float x);
```

In this case, the vector versions of `f` and `g` operate on vectors consisting of 2 and 4 lanes, both with and without an additional lane masking parameter.

For the example, the available (unmasked) signatures associated to `f` and `g` are:

- `float32x2_t _ZGVnN2v_f(float64x2_t vx);` 2-lane `f`;
- `float64x2_t _ZGVnN2v_g(float32x2_t vx);` 2-lane `g`;
- `float32x4_t _ZGVnN4v_f(float64x4_t vx);` 4-lane `f`;
- `float64x4_t _ZGVnN4v_g(float32x4_t vx);` 4-lane `g`;

It is possible to tune the number of lanes using the `simdlen(N)` clause, where $N = 2^k$ for $k \geq 0$. No other values of `simdlen` are allowed.

```
#pragma omp declare simd simdlen(2)
short foo(int64_t x, uint32_t y, int8_t z);
// 2-lane version.
int16x2_t _ZGVnN2vvv_foo(int64x2_t vx, uint32x2_t vy, int8x2_t vz);

#pragma omp declare simd simdlen(4)
short foo(int64_t x, uint32_t y, int8_t z);
// 4-lane version.
int16x4_t _ZGVnN4vvv_foo(int64x4_t vx, uint32x4_t vy, int8x4_t vz);
```

Note

Because AArch64 Advanced SIMD uses the first 8 SIMD registers for passing parameters and returning values, it is recommended that the value passed to `simdlen` is such that the signature of the vector function does not use more than 8 input registers, or more than 8 return registers.

4.7 SVE Examples

In case of the functions `float f(double)`, `double g(float)` and `short foo(int64_t, int32_t, int8_t)`, the use of `#pragma omp declare simd` will generate the following function signatures:

- `svfloat32_t _ZGVsMxv_f(svfloat64_t, svbool_t)` VLA signature for the vector version of `f`;
- `svfloat64_t _ZGVsMxv_g(svfloat32_t, svbool_t)` VLA signature for the vector version of `g`;
- `svint16_t _ZGVsMxvvv_foo(svint64_t, svint32_t, svint8_t, svbool_t)` VLA signature for the vector version of `foo`.

Note that the `svbool_t` parameter is described in [SVE masking](#).

```
// Example with explicit `simdlen` for SVE.

#pragma omp declare simd simdlen(10) notinbranch
#pragma omp declare simd simdlen(16) notinbranch
int32_t foo(int32_t x);

// No 10-lane version generated because ten 4-byte lanes do not
// fit an SVE register.
// SVE 512-bit - widest type is 4 bytes -> 16 lanes
svint32_t _ZGVsM16v_foo(svint32_t vx, svbool_t vmask);

#pragma omp declare simd simdlen(8)
float bar(double x, double y);

// widest type is 8 bytes
// SVE 512-bit -> 8 lanes
svfloat32_t _ZGVsM8vv_bar(svfloat64_t vx, svfloat64_t vy,
                        svbool_t vmask);
```

4.8 Linear parameters examples

Input parameters marked with a `linear` clause need special handling. In particular, the linear clause specifies an implicit vector of values or addresses, depending on the type of the clause.

4.8.1 Linear clause for integral parameters

Meaning of `linear` clause when `x` is an *integral parameter*.

Clause	MAP(<code>x</code>)	Mangled parameter name when <code>s</code> is:		Constraints at lane <code>i</code> of the implicit vector
		Compile time constant	uniform parameter	
<code>linear(x:s)</code>	<code>x</code>	<code>"l" + s</code>	<code>"ls" + pos(s)</code>	<code>x_i = x + i * s</code>
<code>linear(val(x):s)</code>				
<code>linear(uval(x):s)</code>	n/a	n/a	n/a	n/a
<code>linear(ref(x):s)</code>				

4.8.2 Linear clause for pointer parameters

Meaning of `linear` clause when `x` is a *pointer*.

Clause	MAP(<code>x</code>)	Mangled parameter name when <code>s</code> is:		Constraints at lane <code>i</code> of the implicit vector
		Compile time constant	uniform parameter	
<code>linear(x:s)</code>	<code>x</code>	<code>"l" + s * sizeof(*x)</code>	<code>"ls" + pos(s)</code>	<code>x_i = x + i * s</code>
<code>linear(val(x):s)</code>				
<code>linear(uval(x):s)</code>	n/a	n/a	n/a	n/a
<code>linear(ref(x):s)</code>				

4.8.3 Linear clause for integral reference parameters

Meaning of linear clause when x is an integral reference (C++ and Fortran dummy parameters only).

Clause	MAP(x)	Mangled parameter name when s is:		Constraints at lane i of the implicit vector
		Compile time constant	uniform parameter	
linear(x:s)	[&x_0, &x_1, ..., &x_i, ...]	"L" + s	"Ls" + pos(s)	x_i = x + s * i
linear(val(x):s)				
linear(uval(x):s)	x	"U" + s	"Us" + pos(s)	x_i = x + s * i and &x_i = &x
linear(ref(x):s)	x	"R" + s * sizeof(x)	"Rs" + pos(s)	&x_i = &x + s * i

```
// C examples for the ``linear`` clause.

// The same rules apply to dummy arguments passed by value in
// Fortran. Note that the function signatures for the ``val`` modifier
// are the same as when no modifier is present.

// Advanced SIMD
#pragma omp declare simd linear(i)
float bar(int32_t i);
// 2-lane version
float32x2_t _ZGVnN2l_bar(int32_t);
// 4-lane version
float32x4_t _ZGVnN4l_bar(int32_t);

#pragma omp declare simd linear(x)
float foo(double *x);
// 2-lane version
float32x2_t _ZGVnN2l8_foo(double *);
// 4-lane version
float32x4_t _ZGVnN4l8_foo(double *);

// SVE
#pragma omp declare simd linear(i)
float bax(int32_t i);
// VLA version
svfloat32_t _ZGVsMxl_bax(int32_t, svbool_t);

#pragma omp declare simd linear(x)
float bax(double *x);
// VLA version with signature
svfloat32_t _ZGVsMxl8_bax(double *, svbool_t);
```

```
// C++ examples for ``linear`` clause when using reference parameters.

// The same function signature is generated for dummy arguments
// passed by reference in Fortran. For simplicity, the masked version
// for Advanced SIMD is not shown.

#pragma omp declare simd linear(ref(x))
```

```

int32_t g_ref(int32_t &x); // The vector version holds a pointer to x
// Advanced SIMD - 2-lane version
int32x2_t _ZGVnN2R4_g_ref(int32_t *);
// Advanced SIMD - 4-lane version
int32x4_t _ZGVnN4R4_g_ref(int32_t *);
// SVE - VLA version
svint32_t _ZGVsMxR4_g_ref(int32_t *, svbool_t);

#pragma omp declare simd linear(val(x))
int32_t g_val(int32_t &x); // vector of integral values
// Advanced SIMD - 2-lane version
int32x2_t _ZGVnN2L4_g_val(uint64x2_t vxp);
// Advanced SIMD - 4-lane version
int32x4_t _ZGVnN4L4_g_val(uint64x4_t vxp);
// SVE - VLA version
svint32_t _ZGVsMxL4_g_val(svuint64_t vxp , svbool_t);

#pragma omp declare simd linear(uval(x))
int32_t g_uval(int32_t &x); // scalar, used to produce a vector of integral values from x
// Advanced SIMD - 2-lane version
int32x2_t _ZGVsN2U4_g_uval(int32_t *);
// Advanced SIMD - 4-lane version
int32x4_t _ZGVsN4U4_g_uval(int32_t *);
// SVE - VLA version
svint32_t _ZGVsMxU4_g_uval(int32_t *, svbool_t);

```

5 User defined vector functions

Warning

The context of this chapter is at **Beta** level. See [Current status and anticipated changes](#). Any feedback should be provided via the [issue tracker page on GitHub](#).

It is possible to map a scalar function f to a user-defined vector function F by using the directive `#pragma omp declare variant`. This pragma was introduced in version 5.0 of the OpenMP standard.

The following table shows the traits introduced by this Vector Function ABI.

5.1 AArch64 Variant Traits

AArch64 traits for OpenMP contexts.

Trait set	Trait value	Notes
device	<code>isa("simd")</code>	Advanced SIMD call.
device	<code>isa("sve")</code>	SVE call.
device	<code>arch("march-list")</code>	Used to match <code>-march=march-list</code> from the compiler.

The scalar function f that is decorated with a `declare variant` directive with a `simd` trait in the `construct set` is mapped to the vector function F according to the following rules:

1. The signature of F must be the same as that obtained by f when decorated with a `declare simd` directive that matches the `simd` construct specified in the `declare variant` directive, according to the rules specified in [Vector function signature](#).
2. The `device` traits defined in table [AArch64 Variant Traits](#) must be used to narrow the context for matching purposes:
 1. `isa("simd")` targets Advanced SIMD function signatures.
 2. `isa("sve")` targets SVE function signatures.
 3. Either `isa("simd")` or `isa("sve")` must be specified.
 4. The `arch` traits of the `device` set is optional, and it accepts any value that can be passed to the compiler via the command line option `-march`.
3. The `extension("scalable")` trait of the `implementation set` informs the compiler that the `simdlen` clause of the `simd` construct *must* be omitted to target all vector lengths. Its use in a `declare variant` directive is equivalent to having no `simdlen` on `#pragma omp declare simd` when targeting SVE.
4. Using `extension("scalable")` when using `isa("simd")` is invalid.

Note

Decorating a scalar function f with the pragma does not automatically make the vector function F use the vector calling conventions in [Vector Procedure Call Standard](#). The vector function will only use the vector calling conventions if it is marked with the `aarch64_vector_pcs` attribute. The vector function does not need to use the vector calling conventions, although it is recommended in general.

5.2 Examples

```
// User defined `cosine` function for Advanced SIMD.
#pragma omp declare variant(UserCos) \
    match(construct={simd(simdlen(2), notinbranch)}, device={isa("simd")})
double cos(double x);

float64x2_t UserCos(float64x2_t vx);
```

```
// User defined `sincosf` function for VLA SVE.
#pragma omp declare variant(UserSinCos) \
    match(construct={simd(notinbranch, linear(sin, cos))}, \
        device={isa("sve")}, implementation={extension("scalable")})
void sincosf(float in, float *sin, float *cos);

void UserSinCos(svfloat32_t vin, float *sin, float *cos, svbool_t vmask);
```

```
// Advanced SIMD function in an SVE context.
#pragma omp declare variant(F) \
    match(construct={simd(simdlen(4), inbranch)}, \
        device={isa("simd")})
double f(int x);

float64x4_t F(int32x4_t vx, uint32x4_t vmask);
```

```
// VLS version targeting SVE.
#pragma omp declare variant(F) \
    match(construct={simd(simdlen(6), inbranch)}, \
        device={isa("sve")})
double f(int x);

svfloat64_t F(svint32_t vx, svbool_t vmask);
```

```
// Matching via `-march`.
#pragma omp declare variant(H) \
    match(construct={simd(notinbranch)}, \
        implementation={extension("scalable")}, \
        device={isa("sve"), arch("armv8.2-a+sve")})
int h(int x);

svint32_t H(svint32_t vx, svbool_t vmask);
```

```
// Invalid use. This vector signature cannot be derived from the scalar
// function by means of `#pragma omp declare simd`.
#pragma omp declare variant(G) \
```

```
    match(construct={simd(simdlen(2),notinbranch)}, device={isa("sve")})  
char g(double x);  
  
svuint8_t G(float64x2_t vx);
```

6 Masking

The `inbranch` and `notinbranch` clauses define whether or not a vector function should accept a masking parameter.

In all cases, the masking parameter is added to the vector function signature as the last parameter. The following table summarizes the behavior.

Notice that for SVE, masking is present regardless of whether `inbranch` or `notinbranch` is used.¹

Masked signature generation for `[not]inbranch` clause.

	Advanced SIMD		SVE	
	Masked	Unmasked	Masked	Unmasked
<code>#pragma omp declare simd</code>	Yes	Yes	Yes	No
<code>#pragma omp declare simd inbranch</code>	Yes	No	Yes	No
<code>#pragma omp declare simd notinbranch</code>	No	Yes	Yes	No

6.1 Inactive lanes

In a masked vector function, the contents of the inactive lanes of the input parameters and the inactive lanes of the return value are undefined.

6.2 Advanced SIMD masking

For Advanced SIMD, the type of the mask is generated using `uint[NDS(f)*8]_t`-based vectors.

All bits are set to one for active lanes, and all bits are set to zero for inactive lanes.

Note

The narrowest vector input parameter is chosen over the widest one because masking is often intended for lane masking, and not for bit masking of the vector lanes. Using the narrowest vector input parameter also limits the number of parameter registers needed to pass the mask.

Note

Because the masking is done using SIMD data registers, to avoid performance degradation it is recommended that the addition of the mask parameter does not overflow the maximum number of 8 vector input registers.

```
#pragma omp declare simd simdlen(2) inbranch
float f(double);
// 2-lane masked version
float32x2_t _ZGVnM2v_f(float64x2_t, uint32x2_t);

#pragma omp declare simd simdlen(2) inbranch
double g(float);
// 2-lane masked version
float64x2_t _ZGVnM2v_g(float32x2_t, uint32x2_t);
```

```

#pragma omp declare simd inbranch
float f(double); // -> float32x2_t(float64x2_t, uint32x2_t)
// 2 and 4-lane masked version
float32x2_t _ZGVnM2v_f(float64x2_t, uint32x2_t);
float32x4_t _ZGVnM4v_f(float64x4_t, uint32x4_t);

#pragma omp declare simd inbranch
double g(float);
// 2 and 4-lane masked version
float64x2_t _ZGVnM2v_g(float32x2_t, uint32x2_t);
float64x4_t _ZGVnM4v_g(float32x4_t, uint32x4_t);

#pragma omp declare simd simdlen(8) inbranch
float f(double);
// 8-lane masked version
float32x8_t _ZGVnM8v_f(float64x8_t, uint32x8_t);

#pragma omp declare simd simdlen(8) inbranch
double g(float);
// 8-lane masked version
float64x8_t _ZGVnM8v_g(float32x8_t, uint32x8_t);

```

Note

Using a mask parameter in AArch64 Advanced SIMD is not generally recommended for functions that operate on scalars of different widths, as widening of the input mask for wider types might require using call-preserved temporary registers (V8-V23).

6.2.1 Example on Complex Masking

Example of mask parameters for complex values.

```

#pragma omp declare simd inbranch
int32_t foo(_Complex double x);
// Advanced SIMD, 2-lane versions.
// Each logical lane of the mask is a 4 byte sequence,
// either 0x00000000 or 0xffffffff.
int32x2_t _ZGVnM2v_foo(float64x4_t vx, uint32x2_t vmask);

#pragma omp declare simd inbranch
float complex baz(double complex x);
// Double precision complex value -> 16 byte structure
// 2-lane Advanced SIMD.
// The narrowest type is an 8 byte structure, so mask
// is uint64x2_t
float32x4_t _ZGVnM2v_baz(float64x4_t vx, uint64x2_t vmask);

#pragma omp declare simd inbranch
double complex bar(float x, float y);
// Advanced SIMD, 2, and 4-lane.
float64x4_t _ZGVnM2vv_bar(float32x2_t vx, float32x2_t vy, uint32x2_t vmask);
float64x8_t _ZGVnM4vv_bar(float32x4_t vx, float32x4_t vy, uint32x4_t vmask);

```

6.3 SVE masking

For SVE vector functions, whether length-agnostic or length-specific, masked signatures are generated by adding a `svbool_t` mask (or *predicate* in SVE terms) as the last parameter.

```
#pragma omp declare simd
#pragma omp declare simd inbranch
#pragma omp declare simd notinbranch
float f(double);
// SVE - VLA
// Notice that the default behavior is not affected by `inbranch`
// or `notinbranch`.
svfloat32_t _ZGVsMxv_f(svfloat64_t, svbool_t);

#pragma omp declare simd
double g(float);
// SVE - VLA
svfloat64_t _ZGVsMxv_f(svfloat32_t, svbool_t);

#pragma omp declare simd simdlen(4)
float f(double);
// SVE - VLS - > implies a 256-bit implementation
svfloat32_t _ZGVsM4v_f(svfloat64_t, svbool_t);

#pragma omp declare simd simdlen(4)
double g(float);
// SVE - VLS - > implies a 256-bit implementation
svfloat64_t _ZGVsM4v_g(svfloat32_t, svbool_t);
```

6.3.1 Generating the predicate value of the mask parameter

The logical lane subdivision of the predicate corresponds to the lane subdivision of the vector data type generated for the widest data type, with one bit in the predicate lane for each byte of the data lane. Active logical lanes of the predicate have the least significant bit set to 1, and the rest set to zero. The bits of the inactive logical lanes of the predicate are set to zero. This method ensures that:

1. The inactive lanes of unpacked vectors do not get treated erroneously as active (see example `foo`).
2. The correct predicate can be generated programmatically from the input predicate for those types of the scalar signature whose layout requires more than 1 bit per active lane.

In the function `foo` of the following example, the widest data type subdivision selects 8-byte wide lanes. Therefore, the active lanes in the predicate will be represented by the 8-bit sequence `00000001`. The original input predicate works for all the types in the signature but not for the `vy` parameter. The callee must generate a new predicate for it, that carries the bit sequence `00010001` for the active lanes, so that the additional bytes of the logical lane associated to the complex type are correctly marked as active.

```
#pragma omp declare simd
double foo(double x, _Complex float y);

// VLA SVE
svfloat64_t _ZGVsMxv_foo(svfloat64_t vx, svfloat32_t vy,
                        svbool_t vmask);
// vmask active lane value: 00000001
// vy active lane value:    00010001
```

7 Additional examples

Throughout the following examples, for a given function f , we define $NDS(f)$ and $WDS(f)$ as the *Narrowest (and respectively, Widest) Size of f* as the size in bytes of the narrowest (and respectively, the widest) among the input parameter types and the return type of the function signature.

The NDS and WDS values are placed next to the vector signature to explain the choice of the vector length of the function. As a reminder, the former is used to select the vector length when targeting Advanced SIMD vectorization, the latter to select the vector length when targeting VLS SVE functions by using the `simdlen` clause.

```
// Name mangling example for the SIMD directives with no
// decorations.

#pragma omp declare simd
int32_t foo(int32_t x);

// Advanced SIMD - NDS(foo) = 4 -> 2 and 4 lanes
int32x2_t _ZGVnN2v_foo(int32x2_t vx);
int32x2_t _ZGVnM2v_foo(int32x2_t vx, uint32x2_t vmask);
int32x4_t _ZGVnN4v_foo(int32x4_t vx);
int32x4_t _ZGVnM4v_foo(int32x4_t vx, uint32x4_t vmask);

// VLA SVE
svint32_t _ZGVsMxv_foo(svint32_t vx, svbool_t vmask);
```

```
// Example mangling for a function with `uniform` and `linear`
// clause, with `val` modifier. The `inbranch` clause generates only
// the masked version for Advanced SIMD.

#pragma omp declare simd inbranch uniform(x) linear(val(i):4)
int32_t foo(int32_t *x, int32_t i);

// Advanced SIMD - NDS(foo) = 4 -> 2 and 4 lanes
int32x2_t _ZGVnM2ul4_foo(int32_t *x, int32_t i, uint32x2_t vmask);
int32x4_t _ZGVnM4ul4_foo(int32_t *x, int32_t i, uint32x4_t vmask);

// VLA SVE
svint32_t _ZGVsMxul4_foo(int32_t *x, int32_t i, svbool_t vmask);
```

```
// Example of function name mangling when a runtime linear step is
// specified in the `linear` clause.

#pragma omp declare simd inbranch uniform(x,c) linear(i:c)
int32_t foo(int32_t *x, int32_t i, uint8_t c);

// Advanced SIMD - NDS(foo) = 1 -> 8 and 16 lanes
int32x4x2_t _ZGVnM8uls2u_foo(int32_t *x, int32_t i, uint8_t c, uint32x8_t vmask);
int32x4x4_t _ZGVnM16uls2u_foo(int32_t *x, int32_t i, uint8_t c, uint32x16_t vmask);

// VLA SVE
svint32_t _ZGVsMxuls2u_foo(int32_t *x, int32_t i, uint8_t c, svbool_t vmask);
```

```
// Example of vector function name generation from a fixed length
// simd declaration.

#pragma omp declare simd simdlen(4)
```

```
int32_t foo(int32_t x, float y);

// Advanced SIMD - NDS(foo) = 4 -> 4 lanes
int32x4_t _ZGVnN4vv_foo(int32x4_t vx, float32x4_t vy);
int32x4_t _ZGVnM4vv_foo(int32x4_t vx, float32x4_t vy, uint32x4_t vmask);

// SVE 128-bit - WDS(foo) = 4 -> 4 lanes
svint32_t _ZGVsM4vv_foo(svint32_t vx, svfloat32_t vy, svbool_t vmask);
```

```
// Example with output size bigger than input size.

#pragma omp declare simd
double foo(float x)

// Advanced SIMD - NDS(foo) = 4 -> 2 and 4 lanes
float64x2_t _ZGVnN2v_foo(float32x2_t vx);
float64x2_t _ZGVnM2v_foo(float32x2_t vx, uint32x2_t vmask);
float64x4_t _ZGVnN4v_foo(float32x4_t vx);
float64x4_t _ZGVnM4v_foo(float32x4_t vx, uint32x4_t vmask);

// VLA SVE - input in unpacked
svfloat64_t _ZGVsMxv_foo(svfloat32_t vx, svbool_t vmask);
```

Example with explicit alignment.

```
#pragma omp declare simd linear(x) aligned(x:16) simdlen(4)
int32_t foo(int32_t *x, float y);

// Advanced SIMD - NDS(foo) = 4 -> 4 lanes
int32x4_t _ZGVnN4la16v_foo(int32_t *x, float32x4_t vy);
int32x4_t _ZGVnM4la16v_foo(int32_t *x, float32x4_t vy, uint32x4_t vmask);

// SVE 128-bit - WDS(foo) = 4 -> 4 lanes
svint32_t _ZGVsM4la16v_foo(int32_t *x, svfloat32_t vy, svbool_t vmask);
```

7.1 RGB Example

The following example shows how to handle types that do not map directly to integers, floating-point types or complex types. In this specific case, the rules give the following:

1. MTV(P) = true by rule 3 of [Maps To Vector](#).
2. PBV(P) = false by rule 3 of [Pass By Value](#).
3. Because MTV(P) is true, rule 2 of [Parameter and return value mapping](#) applies.
4. Because PBV(P) is false and MTV(P) is true, rule 3 of [Lane Size of a function parameter / return value](#) applies and therefore LS(P) is sizeof(uintptr_t).
5. The vector of pointers to the output values is passed as the first parameter, as specified in rule 3 of [Parameter and return value mapping](#).

```
// Example with generic types. In this case, the rules lead to
// mapping each concurrent object to pointers.

struct S { uint8_t R,G,B; };
```

```

#pragma omp declare simd notinbranch
S DoRGB(S x);

// Advanced SIMD - NDS(DoRGB) = 8 (LP64 data model)
void _ZGVnN2vv_DoRGB(uint64x2_t out, uint64x2_t vx); // 2-lane
// Advanced SIMD - NDS(DoRGB) = 4 (ILP32 data model)
void _ZGVnN2vv_DoRGB(uint32x2_t out, uint32x2_t vx); // 2-lane
void _ZGVnN4vv_DoRGB(uint32x4_t out, uint32x4_t vx); // 4-lane

// VLA SVE - WDS(DoRGB) = 8 (LP64 data model)
void _ZGVsMxvv_DoRGB(svuint64_t out, svint64_t vx, svbool_t vmask);
// VLA SVE - WDS(DoRGB) = 4 (ILP32 data model)
void _ZGVsMxvv_DoRGB(svuint32_t out, svint32_t vx, svbool_t vmask);

```

```

// Example mangling for a function with `uniform` and `linear`
// clause, for corner case values.

#pragma omp declare simd linear(x:y) uniform(y) linear(z) linear(ref(k):-1) notinbranch
uint32_t foo(int32_t x, int32_t y, int32_t z, int32_t &k) {

// Advanced SIMD - NDS(foo) = 4 -> 2 and 4 lanes
uint32x2_t _ZGVnN2lslulRn4_foo(int32_t x, int32_t y, int32_t z, int32_t *k)
uint32x4_t _ZGVnN4lslulRn4_foo(int32_t x, int32_t y, int32_t z, int32_t *k)

// VLA SVE
svuint32_t _ZGVsMxlslulRn4_foo(int32_t x, int32_t y, int32_t z, int32_t *k, svbool_t vmask);

```

```

// Example mangling for default alignment values (assuming LP64).

typedef struct D { double a[2]; } D_ty;

#pragma omp declare simd \
aligned(x) aligned(y) aligned(z) aligned(S) \
linear(x) linear(y) linear(z) linear(S) notinbranch
int32_t foo(int32_t *x, double *y, uint8_t *z, D_ty * S);

// Advanced SIMD - NDS(foo) = 4 -> 2 and 4 lanes (showing only the 2 lanes one)
int32x2_t _ZGVnN2l4a16l8a16l16l16l16_foo(int32_t *x, double *y, uint8_t *z, D_ty * S)

// VLA SVE (VLS would have the same alignment tokens)
svint32_t _ZGVsMxl4a16l8a16l16l16l16_foo(int32_t *x, double *y, uint8_t *z, D_ty * S, svbool_t)

```


8 Footnotes

1

The reason for using predication by default in SVE is to avoid a scalar tail loop when auto-vectorizing loops. The reason for using predication even for `notinbranch` is to avoid the performance degradation that would occur when porting code that uses functions not guarded by conditional branches that could have been marked as `notinbranch`.